

# Distributed Algorithms and Optimization – Parallel Algorithms

Qin Liu

The Chinese University of Hong Kong

# Background

- Parallel and distributed computing is necessary to handle big data

# Background

- Parallel and distributed computing is necessary to handle big data
  - ▶ Parallel: pthread, thread pool pattern, OpenMP

# Background

- Parallel and distributed computing is necessary to handle big data
  - ▶ Parallel: pthread, thread pool pattern, OpenMP
  - ▶ Distributed: MPI, Hadoop, GraphLab, Spark

# Background

- Parallel and distributed computing is necessary to handle big data
  - ▶ Parallel: pthread, thread pool pattern, OpenMP
  - ▶ Distributed: MPI, Hadoop, GraphLab, Spark
- How to design efficient algorithms in parallel and distributed computing?

# Background

- Parallel and distributed computing is necessary to handle big data
  - ▶ Parallel: pthread, thread pool pattern, OpenMP
  - ▶ Distributed: MPI, Hadoop, GraphLab, Spark
- How to design efficient algorithms in parallel and distributed computing?
- Today, start with **parallel computing** on a single machine with **multiple processors** and **shared Random Access Memory**

# Background

- Parallel and distributed computing is necessary to handle big data
  - ▶ Parallel: pthread, thread pool pattern, OpenMP
  - ▶ Distributed: MPI, Hadoop, GraphLab, Spark
- How to design efficient algorithms in parallel and distributed computing?
- Today, start with **parallel computing** on a single machine with **multiple processors** and **shared Random Access Memory**
  - ▶ Distributed computing involves network communication overhead

# Why focus on parallel algorithms?

- The growth CPU clock-speed slowed down after it has reached around 3 Gigahertz



# Why focus on parallel algorithms?

- The growth CPU clock-speed slowed down after it has reached around 3 Gigahertz
- Now, machines have many cores

# Why focus on parallel algorithms?

- The growth CPU clock-speed slowed down after it has reached around 3 Gigahertz
- Now, machines have many cores
- Because computation has shifted from sequential to parallel, our algorithms have to change as well...

# Outline

Fundamentals of Parallel Algorithm Analysis

The Master Theorem

Matrix Multiplication: Strassen's algorithm

Mergesort

# Fundamentals of Parallel Algorithm Analysis

# Efficiency of Parallel Algorithms

- Sequential RAM Model: one processor and one memory module

# Efficiency of Parallel Algorithms

- Sequential RAM Model: one processor and one memory module
  - ▶ use the number of operations to estimate the wall-clock compute time: e.g., quick sort takes  $O(n \log n)$ , bubble sort takes  $O(n^2)$

# Efficiency of Parallel Algorithms

- Sequential RAM Model: one processor and one memory module
  - ▶ use the number of operations to estimate the wall-clock compute time: e.g., quick sort takes  $O(n \log n)$ , bubble sort takes  $O(n^2)$
- Parallel RAM Model: multiple processors interact with the memory module(s)

# Efficiency of Parallel Algorithms

- Sequential RAM Model: one processor and one memory module
  - ▶ use the number of operations to estimate the wall-clock compute time: e.g., quick sort takes  $O(n \log n)$ , bubble sort takes  $O(n^2)$
- Parallel RAM Model: multiple processors interact with the memory module(s)
  - ▶ usually suppose concurrent read and exclusive write



# Caveats of PRAM



Many variants for an algorithm designer to consider:

- Resolving concurrent writes: Garbage, Arbitrary, Priority, Combination

# Caveats of PRAM

Many variants for an algorithm designer to consider:



- Resolving concurrent writes: Garbage, Arbitrary, Priority, Combination
- Memory speed



# Caveats of PRAM

Many variants for an algorithm designer to consider:

- Resolving concurrent writes: Garbage, Arbitrary, Priority, Combination
- Memory speed



- ▶ layers of cache in between each processor and memory module which all have different read-write speeds

# Notations

Let

$T_1$  = amount of (wall-clock) time algorithm takes on one processor

$T_p$  = amount of (wall-clock) time algorithm takes on  $p$  processors

# Notations

Let

$T_1$  = amount of (wall-clock) time algorithm takes on one processor

$T_p$  = amount of (wall-clock) time algorithm takes on  $p$  processors

Lower bound of  $T_p$ :

$$\frac{T_1}{p} \leq T_p$$

# Work and Depth

## **Work:**

- Defined as the number of operations required

# Work and Depth

## **Work:**

- Defined as the number of operations required
- Enough for sequential models, but no longer holds in a parallel model

# Work and Depth

## **Work:**

- Defined as the number of operations required
- Enough for sequential models, but no longer holds in a parallel model

## **Depth:**

- What if we have infinitude of processors? Would the compute time for an algorithm then be zero?



# Work and Depth

## Work:

- Defined as the number of operations required
- Enough for sequential models, but no longer holds in a parallel model

## Depth:

- What if we have infinitude of processors? Would the compute time for an algorithm then be zero?
- No, because algorithms usually have an **inherently** sequential component to them

# Work and Depth

## Work:

- Defined as the number of operations required
- Enough for sequential models, but no longer holds in a parallel model

## Depth:

- What if we have infinitude of processors? Would the compute time for an algorithm then be zero?
- No, because algorithms usually have an **inherently** sequential component to them
- Known as **depth**

# Representing Algorithms as DAG's




Figure: Example DAG

Represent the dependencies between operations in an algorithm using a directed acyclic graph (DAG)

- each fundamental unit of computation is represented by a node

# Representing Algorithms as DAG's




Figure: Example DAG

Represent the dependencies between operations in an algorithm using a directed acyclic graph (DAG)

- each fundamental unit of computation is represented by a node
- an edge from node  $u$  to node  $v$  means if computation of  $v$  is required as an input to computation  $u$

# Representing Algorithms as DAG's




Figure: Example DAG

Represent the dependencies between operations in an algorithm using a directed acyclic graph (DAG)

- each fundamental unit of computation is represented by a node
- an edge from node  $u$  to node  $v$  means if computation of  $v$  is required as an input to computation  $u$
- it's actually a tree

# Representing Algorithms as DAG's




Figure: Example DAG

Represent the dependencies between operations in an algorithm using a directed acyclic graph (DAG)

- each fundamental unit of computation is represented by a node
- an edge from node  $u$  to node  $v$  means if computation of  $v$  is required as an input to computation  $u$
- it's actually a tree
- the root represents the output of our algorithm

## Representing Algorithms as DAG's (cont'd)

Define **work** to be

$$T_1 = \text{number of nodes in DAG}$$

## Representing Algorithms as DAG's (cont'd)

Define **work** to be

$$T_1 = \text{number of nodes in DAG}$$

With an infinitude of processors, the compute time is given by the depth of the tree

$$T_\infty = \text{depth of computation DAG}$$



# Brent's theorem

## Theorem (Brent's theorem)

*We claim*

$$\frac{T_1}{p} \leq T_p \leq \frac{T_1}{p} + T_\infty$$

The depth  $T_\infty$  measures how far off our algorithm performs relative to the best possible version or how **parallel** an algorithm is.

# Proof of Brent's theorem

**Proof:** On level  $i$  of our DAG, there are  $m_i$  operations. Hence by that definition, since  $T_1$  is the total work of our algorithm,

$$T_1 = \sum_{i=1}^n m_i$$

where we denote  $T_\infty = n$ . For each level  $i$  of the DAG, the time taken by  $p$  processors is given as

$$T_p^i = \left\lceil \frac{m_i}{p} \right\rceil \leq \frac{m_i}{p} + 1.$$

This equality follows from the fact that there are  $m_i$  constant-time operations to be performed at  $m_i$ , and once all lower levels have been completed, these operations share no inter-dependencies. Thus, we may distribute or assign operations uniformly to our processors. The ceiling follows from the fact that if the number of processors not divisible by  $p$ , we require exactly one wall-clock cycle where some but not all processors are used in parallel. Then,

$$T_p = \sum_{i=1}^n T_p^i \leq \sum_{i=1}^n \left( \frac{m_i}{p} + 1 \right) = \frac{T_1}{p} + T_\infty$$

■

# Speed-up

$T_{p,n}$ : the run-time on  $p$  processors given an input of size  $n$ .  
The speed-up of a parallel algorithm:  $\text{SpeedUp}(p, n) = \frac{T_{1,n}}{T_{p,n}}$

## Definition

*If  $\text{SpeedUp}(p, n) = \Theta(p)$ , we say the algorithm is strongly scalable.*

## Definition

*If  $\text{SpeedUp}(p, np) = \frac{T_{1,n}}{T_{p,np}} = \Omega(1)$ , we say the algorithm is weakly scalable.*

# Embarrassingly Parallel



Figure: An Embarrassingly Parallel DAG

- No dependency between operations
- Scalable in the most trivial sense

## Example: Summation

```
1  $s \leftarrow 0$  for  $i \leftarrow 1, 2, \dots, n$  do  
2   |  $s \leftarrow s + a[i]$   
3 end  
4 return  $s$ 
```

Figure: Sequential Summation

$$T_1 = T_2 = \dots = T_\infty = n$$

How to redesign the algorithm?

# Parallel Summation

Instead of

$$((a_1 + a_2) + a_3) + a_4$$

We assign each processor a pair of elements

$$(a_1 + a_2) + (a_3 + a_4)$$

# Parallel Summation

Instead of

$$((a_1 + a_2) + a_3) + a_4$$

We assign each processor a pair of elements

$$(a_1 + a_2) + (a_3 + a_4)$$

Results in

$$T_{\infty} = \log_2 n$$

Hence by Brent's theorem

$$T_p \leq \frac{n}{p} + \log_2 n$$

## Example: Matrix Multiplies

Given two  $n \times n$  matrices,  $A$  and  $B$ , output one  $n \times n$  matrix  $C = AB$ , where

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$



## Example: Matrix Multiplies

Given two  $n \times n$  matrices,  $A$  and  $B$ , output one  $n \times n$  matrix  $C = AB$ , where

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

Work is

$$T_1 = O(n^3)$$

## Example: Matrix Multiplies

Given two  $n \times n$  matrices,  $A$  and  $B$ , output one  $n \times n$  matrix  $C = AB$ , where

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

Work is

$$T_1 = O(n^3)$$

Parallel algorithm:

- Each  $c_{ij}$  takes  $O(n)$  independently
- $T_p \leq O(\frac{n^3}{p} + n)$
- Parallelize each  $c_{ij}$  summation
  - ▶  $T_p \leq O(\frac{n^3}{p} + \log n)$

# The Master Theorem

# Recursive Algorithms

- Easy to parallelize
- The recursive calls don't depend on each other, hence we may assign each to a processor
- How to analyze the time complexity?

# Unrolling the Simplest Case

Suppose the recurrence relation is of the form:

$$T(n) = bT(n/b) + n,$$

where  $b$  is an integer greater than 1 (e.g.,  $b = 2$  for quick sort).

$T(n) = bT\left(\frac{n}{b}\right) + n$	by our definition of $T(n)$
$= b\left[bT\left(\frac{n}{b^2}\right) + \frac{n}{b}\right] + n$	unrolling to the second-level of our tree
$= b^2T\left(\frac{n}{b^2}\right) + b \cdot \frac{n}{b} + n$	re-arranging
$= b^2\left[bT\left(\frac{n}{b^3}\right) + \frac{n}{b^2}\right] + b \cdot \frac{n}{b} + n$	unrolling to the third level of our tree
$= b^3T\left(\frac{n}{b^3}\right) + b^2 \cdot \frac{n}{b^2} + b \cdot \frac{n}{b} + n$	re-arranging
$= \dots$	
$= n \log_b(n)$	

# General Form of the Master Theorem

Suppose  $T(n) = aT(\frac{n}{b}) + f(n)$

1. If  $f(n) \in O(n^c)$  where  $c < \log_b a$ , then

$$T(n) \in \Theta(n^{\log_b a})$$

2. If for some constant  $k \geq 0$ ,  $f(n) \in \Theta(n^c \log^k n)$  where  $c = \log_b a$ , then

$$T(n) \in \Theta(n^c \log_b^{k+1} n)$$

3. If  $f(n) \in O(n^c)$  where  $c > \log_b a$  and also

$$af(\frac{n}{b}) \leq kf(n)$$

for some constant  $k < 1$ , then

$$T(n) \in \Theta(f(n))$$

# Examples

- $T(n) = T(n/b) + 1.$

# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$



# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$
  - ▶  $T(n) = \log_b n$

# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$
  - ▶  $T(n) = \log_b n$
  - ▶ E.g., binary search  $O(\log_2 n)$

# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$
  - ▶  $T(n) = \log_b n$
  - ▶ E.g., binary search  $O(\log_2 n)$
- $T(n) = T(n/b) + n$ .

# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$
  - ▶  $T(n) = \log_b n$
  - ▶ E.g., binary search  $O(\log_2 n)$
- $T(n) = T(n/b) + n$ .
  - ▶ Applying case 3,  $c = 1, k \in (1/b, 1)$

# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$
  - ▶  $T(n) = \log_b n$
  - ▶ E.g., binary search  $O(\log_2 n)$
- $T(n) = T(n/b) + n$ .
  - ▶ Applying case 3,  $c = 1, k \in (1/b, 1)$
  - ▶  $T(n) = \Theta(n)$

# Examples

- $T(n) = T(n/b) + 1$ .
  - ▶ Applying case 2,  $c = 0, k = 0$
  - ▶  $T(n) = \log_b n$
  - ▶ E.g., binary search  $O(\log_2 n)$
- $T(n) = T(n/b) + n$ .
  - ▶ Applying case 3,  $c = 1, k \in (1/b, 1)$
  - ▶  $T(n) = \Theta(n)$
  - ▶ E.g., quick select  $O(n)$

## General Recursive Relations

Example:  $T(n) = T(\sqrt{n}) + 1$ , with a base case  $T(2) = 1$ .  
We have to unroll the recurrence by hand to find a solution.

$T(n) = T(n^{1/2}) + 1$	our definition of $T(n)$
$= (T(n^{1/4}) + 1) + 1$	unrolling to second level
$= ((T(n^{1/8}) + 1) + 1) + 1$	unrolling to third level...

## General Recursive Relations

Example:  $T(n) = T(\sqrt{n}) + 1$ , with a base case  $T(2) = 1$ .  
We have to unroll the recurrence by hand to find a solution.

$T(n) = T(n^{1/2}) + 1$	our definition of $T(n)$
$= (T(n^{1/4}) + 1) + 1$	unrolling to second level
$= ((T(n^{1/8}) + 1) + 1) + 1$	unrolling to third level...

Let  $k$  denote the number of recursive calls to reach base case:

$$n^{1/2^k} = 2$$

$$\frac{1}{2^k} \log_2 n = \log_2 2$$

$$k = \log_2 \log_2 n$$



## Example: All Prefix Sum

Given

$$A = [a_1, a_2, \dots, a_n],$$

output

$$R = [r_0, r_1, r_2, \dots, r_n],$$

where  $r_k = \sum_{i=1}^k a_i$  and  $r_0 = 0$ .

### Algorithm 2: Prefix Sum

**Input:** All prefix sum for an array  $A$

- 1 **if** *size of  $A$  is 1* **then**
- 2     **return** *only element of  $A$*
- 3 **end**
- 4 Let  $A'$  be the sum of adjacent pairs
- 5 Compute  $R' = \text{AllPrefixSum}(A')$  // Note:  $R'$  has every other element of  $R$
- 6 Fill in missing entries of  $R'$  using another  $\frac{n}{2}$  processors

## Example: All Prefix Sum (cont'd)

Given

$$A = [a_1, a_2, \dots, a_n],$$

output

$$R = [r_0, r_1, r_2, \dots, r_n],$$

where  $r_k = \sum_{i=1}^k a_i$ .

## Example: All Prefix Sum (cont'd)

Given

$$A = [a_1, a_2, \dots, a_n],$$

output

$$R = [r_0, r_1, r_2, \dots, r_n],$$

where  $r_k = \sum_{i=1}^k a_i$ .

Let

$$A' = [a_1 + a_2, a_3 + a_4, \dots, a_{n-1} + a_n]$$

$$R' = [r_2, r_4, \dots, r_n]$$

For  $i$  is odd

$$r_i = r_{i-1} + a_i$$

# Algorithm Analysis

Let  $T_1 = W(n)$  and  $T_\infty = D(n)$ , we have

$$W(n) = W(n/2) + O(n) = O(n)$$

$$D(n) = D(n/2) + O(1) = O(\log n)$$

So

$$T_p \leq O(n/p + \log n)$$

# Matrix Multiplication: Strassen's algorithm

# Matrix Multiplies

- Naive way:  $O(n^3)$
- Strassen's sequential algorithm:  $O(n^{\log_2 7}) = O(n^{2.81})$
- Idea – Block Matrix Multiplication

$$\mathbf{C} = \mathbf{AB} \quad \mathbf{A}, \mathbf{B}, \mathbf{C} \in R^{2^n \times 2^n}$$

$$\mathbf{A} = \begin{bmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{bmatrix}, \mathbf{B} = \begin{bmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{bmatrix}, \mathbf{C} = \begin{bmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{bmatrix}$$

# Block Matrix Multiplication

$$\mathbf{C}_{1,1} = \mathbf{A}_{1,1}\mathbf{B}_{1,1} + \mathbf{A}_{1,2}\mathbf{B}_{2,1}$$

$$\mathbf{C}_{1,2} = \mathbf{A}_{1,1}\mathbf{B}_{1,2} + \mathbf{A}_{1,2}\mathbf{B}_{2,2}$$

$$\mathbf{C}_{2,1} = \mathbf{A}_{2,1}\mathbf{B}_{1,1} + \mathbf{A}_{2,2}\mathbf{B}_{2,1}$$

$$\mathbf{C}_{2,2} = \mathbf{A}_{2,1}\mathbf{B}_{1,2} + \mathbf{A}_{2,2}\mathbf{B}_{2,2}$$

## Parallelizing the Algorithm

8 matrix multiplies between matrices of size  $n/2 \times n/2$  and 4 matrix additions:

$$W(n) = 8W(n/2) + O(n^2)$$

By the Master Theorem,

$$W(n) = O(n^3)$$

No progress compared to the naive way



# Strassen's Algorithm

Reduce the number of sub-calls to matrix-multiplies to 7:

$$\mathbf{M}_1 := (\mathbf{A}_{1,1} + \mathbf{A}_{2,2})(\mathbf{B}_{1,1} + \mathbf{B}_{2,2})$$

$$\mathbf{M}_2 := (\mathbf{A}_{2,1} + \mathbf{A}_{2,2})\mathbf{B}_{1,1}$$

$$\mathbf{M}_3 := \mathbf{A}_{1,1}(\mathbf{B}_{1,2} - \mathbf{B}_{2,2})$$

$$\mathbf{M}_4 := \mathbf{A}_{2,2}(\mathbf{B}_{2,1} - \mathbf{B}_{1,1})$$

$$\mathbf{M}_5 := (\mathbf{A}_{1,1} + \mathbf{A}_{1,2})\mathbf{B}_{2,2}$$

$$\mathbf{M}_6 := (\mathbf{A}_{2,1} - \mathbf{A}_{1,1})(\mathbf{B}_{1,1} + \mathbf{B}_{1,2})$$

$$\mathbf{M}_7 := (\mathbf{A}_{1,2} - \mathbf{A}_{2,2})(\mathbf{B}_{2,1} + \mathbf{B}_{2,2})$$

## Strassen's Algorithm (cont'd)

$$\mathbf{C}_{1,1} = \mathbf{M}_1 + \mathbf{M}_4 - \mathbf{M}_5 + \mathbf{M}_7$$

$$\mathbf{C}_{1,2} = \mathbf{M}_3 + \mathbf{M}_5$$

$$\mathbf{C}_{2,1} = \mathbf{M}_2 + \mathbf{M}_4$$

$$\mathbf{C}_{2,2} = \mathbf{M}_1 - \mathbf{M}_2 + \mathbf{M}_3 + \mathbf{M}_6$$

## Strassen's Algorithm (cont'd)

Work:

$$W(n) = 7W(n/2) + O(n^2)$$

By the Master Theorem,

$$W(n) = O(n^{\log_2 7})$$

## Strassen's Algorithm (cont'd)

Work:

$$W(n) = 7W(n/2) + O(n^2)$$

By the Master Theorem,

$$W(n) = O(n^{\log_2 7})$$

Since the sub-calls can be parallelized and the depth of matrix-additions is  $O(1)$

$$D(n) = D(n/2) + O(1) = O(\log n)$$

## Strassen's Algorithm (cont'd)

Work:

$$W(n) = 7W(n/2) + O(n^2)$$

By the Master Theorem,

$$W(n) = O(n^{\log_2 7})$$

Since the sub-calls can be parallelized and the depth of matrix-additions is  $O(1)$

$$D(n) = D(n/2) + O(1) = O(\log n)$$

By Brent's theorem

$$T_p \leq O\left(\frac{n^{2.81}}{p} + \log n\right)$$

# Drawbacks of Divide and Conquer

## Communication cost:

- Parallel RAM model assumes zero communication costs between processors
- Realistically, we never have efficient communication
- In clusters of computers, Strassen's algorithm is impractical
- Need to chop up matrices which incurs lots of shuffle cost

## Drawbacks of Divide and Conquer (cont'd)

Big  $\mathcal{O}$  and big constants:

- In Strassen's algorithm, the  $O(n^2)$  term requires  $20 \cdot n$  operations
- When the data is distributed across machines, we can only afford to pass it one time
- Big  $\mathcal{O}$  notion is good to get started and throw away super bad algorithms
- Sometimes we need to look closely

# Mergesort



# Mergesort

## Algorithm 1: Merge Sort

**Input** : Array  $A$  with  $n$  elements

**Output:** Sorted  $A$

```
1  $n \leftarrow |A|$ 
2 if  $n$  is 1 then
3   return  $A$ 
4 end
5 else
6   // (In Parallel)
7    $L \leftarrow \text{MERGESORT}(A[0, \dots, n/2])$  // Indices  $0, 1, \dots, \frac{n}{2} - 1$ 
8    $R \leftarrow \text{MERGESORT}(A[n/2, \dots, n])$  // Indices  $\frac{n}{2}, \frac{n}{2} + 1, \dots, n - 1$ 
9   return  $\text{MERGE}(L, R)$ 
9 end
```

# Mergesort (cont'd)

## Algorithm 2: Merge

**Input** : Two sorted arrays  $A, B$  each of length  $n$

**Output**: Merged array  $C$ , consisting of elements of  $A$  and  $B$  in sorted order

```
1  $a \leftarrow$  pointer to head of array  $A$  (i.e. pointer to smallest element in  $A$ )
2  $b \leftarrow$  pointer to head of array  $B$  (i.e. pointer to smallest element in  $B$ )
3 while  $a, b$  are not null do
4   Compare the value of the element at  $a$  with the value of the element at  $b$ 
5   if  $value(a) < value(b)$  then
6     add value of  $a$  to output  $C$ 
7     increment pointer  $a$  to next element in  $A$ 
8   end
9   else
10    add value of  $b$  to output  $C$ 
11    increment pointer  $b$  to next element in  $B$ 
12  end
13 end
14 if elements remaining in either  $a$  or (exclusive)  $b$  then
15   Append these sorted elements to our sorted output  $C$ 
16 end
17 return  $C$ 
```

# Naive Parallelization

We have

$$W(n) = 2W(n/2) + O(n) = O(n \log n)$$

$$D(n) = D(n/2) + O(n) = O(n)$$

# Naive Parallelization

We have

$$W(n) = 2W(n/2) + O(n) = O(n \log n)$$

$$D(n) = D(n/2) + O(n) = O(n)$$

Using Brent's theorem

$$T_p \leq O\left(\frac{n \log n}{p} + n\right)$$

# Naive Parallelization

We have

$$W(n) = 2W(n/2) + O(n) = O(n \log n)$$

$$D(n) = D(n/2) + O(n) = O(n)$$

Using Brent's theorem

$$T_p \leq O\left(\frac{n \log n}{p} + n\right)$$

Not much speed-up compared to sequential version.  
Bottleneck lies in merge.

# Improved Parallelization

Merge two sorted array  $A$  and  $B$  in parallel. Suppose the output is  $C$ , for any element  $x$  in  $A$  or  $B$

$$\text{rank}_C(x) = \text{rank}_A(x) + \text{rank}_B(x)$$

**Algorithm 3:** Parallel Merge

**Input** : Two sorted arrays  $A, B$  each of length  $n$

**Output:** Merged array  $C$ , consisting of elements of  $A$  and  $B$  in sorted order

1 **for** each  $a \in A$  **do**

2     Do a binary search to find where  $a$  would be added into  $B$ ,

3     The final rank of  $a$  given by  $\text{rank}_C(a) = \text{rank}_A(a) + \text{rank}_B(a)$ .

4 **end**

$n$  parallel binary searches, each takes  $O(\log n/2) = O(\log n)$  time.

## Improved Parallelization (cont'd)

We have

$$W(n) = 2W(n/2) + O(n \log n) = O(n \log^2 n)$$

$$D(n) = D(n/2) + O(\log n) = O(\log^2 n)$$

By Brent's theorem

$$T_p \leq O\left(\frac{n \log^2 n}{p} + \log^2 n\right)$$

Best known algorithm by Richard Cole:

$$T_p \leq O\left(\frac{n \log n}{p} + \log n\right)$$

# References

- CME 323: Distributed Algorithms and Optimization by Reza Zadeh (Stanford)
- Parallel Algorithms by Guy E. Blelloch and Bruce M. Maggs
- Introduction to Algorithms by Cormen, Leiserson, Rivest, Stein